TDS Category Check: Take-Home Brief

Nov 3, 2026 · Savar Verma

Indian companies choose a TDS category for each vendor payment, often without recording why. This design checks those choices before each quarterly return, cites its evidence, and leaves the decision with a person.

Area explored and what surprised me

I explored tax deducted at source (TDS) in India. It makes the payer collect part of the payee's income tax upfront, then deposit and report it. I covered non-salary payments to Indian residents under the Income Tax Act 2025.

Three things surprised me:

- **Calculating TDS is the easy part.** Thresholds, PAN checks and deadlines are ordinary software problems. The judgment sits in one step: deciding what a payment actually is.
- **A correct company can still lose years.** A co-operative bank deducted 2% on connectivity charges, and the officer demanded 10%. The ₹21.15 lakh demand survived first appeal and was deleted only at the tribunal, in a September 2026 report on an appeal numbered from 2023.
- **The cost is moderate, not catastrophic.** Early reading suggested errors would be expensive because the expense itself gets disallowed. Recent rulings lean against that when some tax was deducted. So I scoped the claim down.

Problem statement

Accounts teams can't see when a TDS category has gone wrong, and often can't later explain why they chose it. I believe the category is often set once per vendor and then reused, even when the work changes. This reuse is a hypothesis I have not yet confirmed with a practitioner.

- **Who:** accounts teams at mid-sized Indian companies. I assume large firms have tax teams and small firms rely on an outside CA.
- **Task:** assign each vendor payment a category, which sets the rate and payment code. For example, code 1026 is technical services at 2%, and code 1027 is professional services at 10%.
- **How often:** on every vendor payment, with returns filed quarterly. The error rate is unknown, and I found no public figure.
- **Why now:** since 1 April 2026, the old 194-series labels have been replaced by section 393 and numeric payment codes. Every vendor mapping had to be redone, so this is a natural moment to review them.

What it costs today. These are real disputes, taken from secondary reports:

Reported	Payer	Dispute	Cost and outcome
Sep 2026	H.P. State Co-operative Bank	Connectivity charges: 2% deducted, 10% demanded	₹21.15 lakh demand, deleted at the tribunal
Aug 2026	Not named in the report	Maintenance charges: 2% as contract work, 10% as rent demanded	₹3.43 lakh demand, including ₹1.52 lakh interest; sent back for a fresh decision
Feb 2025	BSNL	Services under the wrong section, for assessment year 2008-09	₹1.02 crore shortfall and ₹5.22 lakh interest

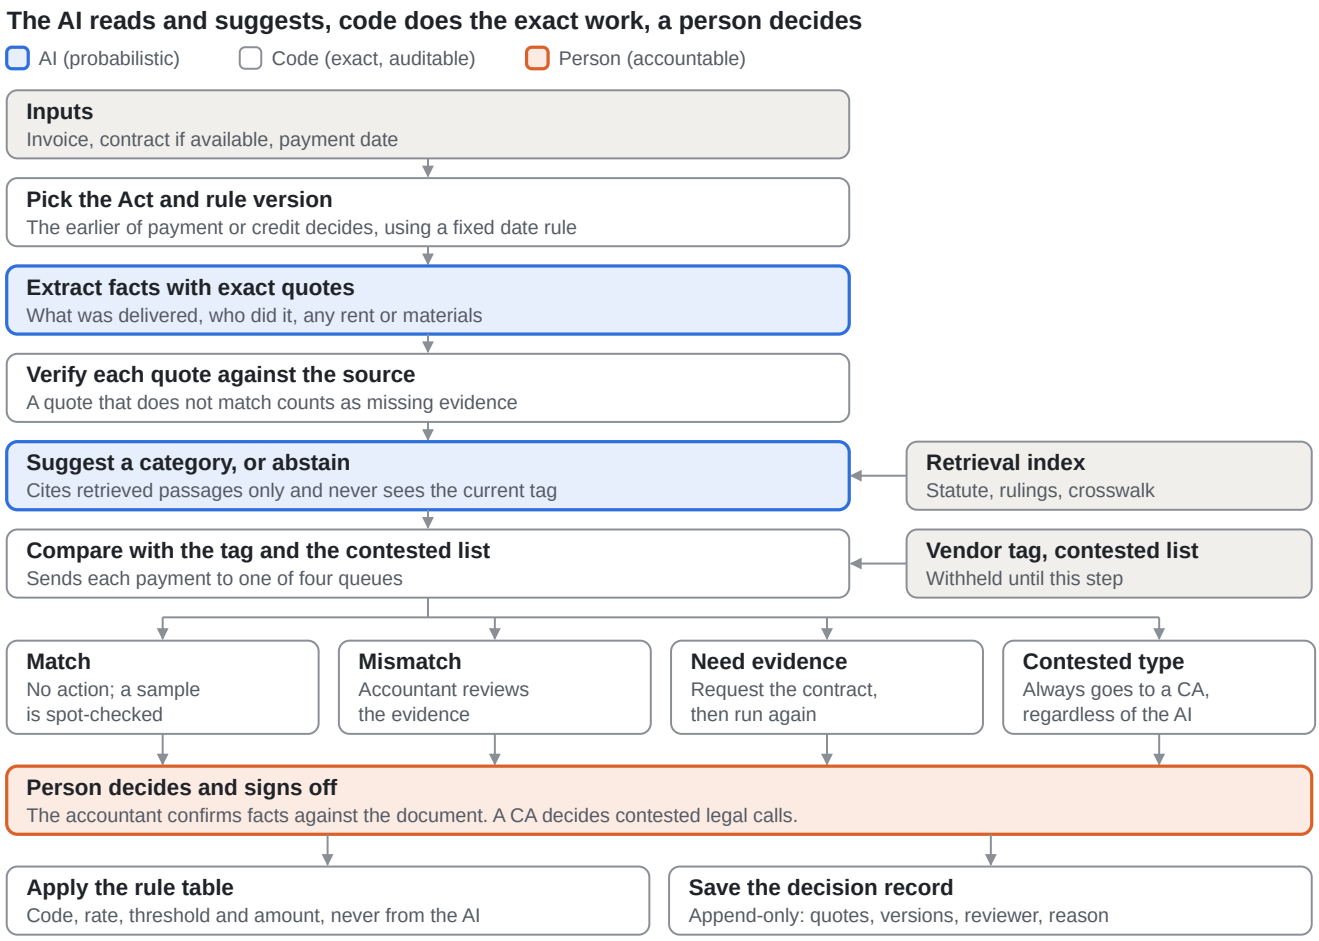
These only show disputes that reached a tribunal, so they don't show frequency. I treat this as a real, recurring and moderately costly problem with an unknown error rate. The added costs are appeal time and the lack of a record to defend the call.

Out of scope: salary, payments to non-residents, TCS, filing returns, depositing tax, and deciding the law on contested questions.

Solution design

Before each quarterly return, the tool re-checks that quarter's vendor payments against their current category. It questions a person's decision instead of replacing it, which makes it easier to trust and to pilot.

- **What goes in:** the invoice, the contract if available, the payment date and the vendor's current tag.
- **What comes out:** a queue for each payment, a suggested category with quoted evidence, the resulting code and amount, and a saved record.
- **Uncertainty and mistakes:** missing evidence means the tool abstains and asks for the contract. It doesn't guess. A person can overrule any suggestion, and the reason is saved. Sampled matches catch errors the tool missed.



Design decisions

Why AI is necessary. The deciding facts sit in free text, and the same words can point to different categories. One tribunal looked for a human element in connectivity charges. In a hospital equipment dispute, the first appeal separated routine maintenance contracts (contract work) from specialised ones (technical services), and the Bombay High Court later ordered each contract to be examined. Keyword rules can't make that distinction, and I know of no public labelled data to train a classifier.

Why the AI is kept narrow. Once the category is known, everything else is a lookup. So the AI interprets text, and code does everything that must be exact.

Part	Handled by	Why
Which Act applies (earlier of payment or credit)	Code	A fixed date rule
Code, rate, threshold, tax amount	Code	Must be exact and auditable
Quotes and citations	Code checks them	A quote must match the source text, and a citation must be a retrieved passage
Facts in the invoice or contract	LLM	Free text
Suggested category	LLM	Judgment, so it is treated as probabilistic
Final decision	Person	Legal accountability

Blind classification. The LLM never sees the vendor's current tag. It classifies from the documents, then code compares the result with the tag. Showing the tag would anchor the model and hide the errors this tool exists to catch.

Routing doesn't trust model confidence. A human-maintained list of contested payment types, seeded from case law, always goes to a CA, however sure the model is. An LLM's confidence isn't reliable evidence that an open legal question is settled.

Where the knowledge lives, and what happens when it expires:

- **Rates and codes** live in code, versioned by effective date. A new version applies only to payments dated after it starts.
- **Statute and rulings** live in a retrieval index, tagged with date and Act.
- **The old-to-new crosswalk** is a reviewed table. Most rulings cite 194C or 194J while filings now use section 393, so plain text search can miss the link.
- **Contested types** live with a human reviewer, because deciding what is unsettled is itself a judgment.
- **The model** is never the authority for a rule. Anything it can't cite is discarded.

- **When a rule or ruling changes**, past decisions that relied on it are flagged for review and are not silently rewritten. For example, one source reports that Finance Act 2026 now treats supply of manpower as contract work, which would take that type off the contested list.

Traceability. Each decision saves an append-only record: document hashes, quoted text, passages used, rule version, model version, reviewer and reason. An auditor can replay how any category was reached. This is also the record the company lacks today when an officer asks why it chose 2%.

Accountability. The tool gives no tax advice to outsiders and files nothing. The company remains responsible as deductor, a named person approves each change, and contested calls go to a qualified CA.

Alternatives rejected

Alternative	Why I rejected it
Chat assistant that accountants prompt (my first idea)	It only helps when someone already suspects a problem, and nobody asks about errors they can't see
One LLM call	Rates and citations could be invented, and nothing can be checked or replayed
Rules only	Rules can't read what the work was from free text
Trained classifier	I know of no public labelled data, and it can't cite why it chose a category
Several cooperating agents	The steps run in a straight line, so more agents add failure points but no knowledge
Automating filing too	Existing software appears to cover this, and it adds liability without solving the judgment problem
Checking at every payment	It catches errors earlier but must sit inside the payment flow. I chose a review before each quarterly return because it is smaller to pilot, accepting that an error can run for up to a quarter

Risks

Risk	How the design handles it
Wrong answers	Blind check, verified quotes, human approval. I would measure wrong-and-unflagged cases and the abstention rate, because overall accuracy hides the failures that matter
Outdated rules	Effective-dated rule table, and flagged re-review when something changes
Unsettled law	An April 2026 article argues that old wrong-section protection may not carry over to a wrong code within section 393, so the 30% disallowance (reported as now in section 35(b)) could apply. It reported no ruling at the time. So old rulings are shown as context, not as settled answers
Confidential data	Vendor contracts and PANs are sensitive. A real deployment needs a private or zero-retention model and masked PANs. This brief uses only public and synthetic data
Liability	The tool recommends and never files. A named person approves each change, and the company stays accountable
Weak reviewers	If the tool lets a less qualified accountant do this work, that person is more likely to accept a wrong suggestion. So they only confirm facts against the document, and contested legal calls go to a CA

Weakest point. The tool only sees documents, and those can misdescribe the real work. No model can detect that.

What would break it. If most invoices are too vague and contracts aren't available, nearly everything lands in the "need evidence" queue. The tool would then create work instead of removing it. I would test this first.

Before spending more time

I have not spoken to a practitioner or built the optional prototype. So the core hypothesis is untested, and these checks come first:

1. **Ask a CA or accounts team:** is the category set per vendor or per invoice? When did a wrong one last cost money? What do you show an officer? Does your software or auditor already catch mismatches? If auditors already catch this reliably, the problem shrinks.
2. **Check the law directly:** my rates and codes come from secondary sources, so they need checking against section 393 and the Income-tax Rules 2026.
3. **Run a small test:** use about 20 synthetic invoices, including vague and contested ones, and measure how many the tool can't decide.

Another jurisdiction. The rule table, crosswalk and rulings would change, while the pipeline would stay. I expect similar value where business payments are widely withheld, such as the Philippines or Indonesia. In the US, where most freelancer payments are reported rather than withheld, this problem would mostly disappear. I have not verified these country details.

Sources

All are secondary sources, and I have not read the orders themselves. The bank report is partly paywalled, so its appeal details and the maintenance-charges case come from search excerpts only.

- Taxscan: H.P. State Co-operative Bank
- TaxCorp: maintenance charges dispute
- TaxHeal: BSNL order
- ABCAUS: hospital maintenance contracts, Bombay High Court
- TaxGuru: Shree Laminators, short deduction
- TaxGuru: wrong code under the 2025 Act
- Tax Garden: section 393 code mapping

Appendix: what I added beyond the AI

Where AI helped. I used Claude to learn TDS from scratch, find rulings, pressure-test the problem, and draft this brief's text from our conversation. Key prompts:

- "Let's solve this. What do I need to know and what do we do with this?"
- "Are we just trying to ease the work of the accountants at a company? ... Or is there a deeper problem we are trying to fix?"
- "i want to personally ask, when is the last time second happened irl?"

What I rejected or corrected.

- **The size of the problem.** Claude's early framing stressed the risk of disallowing the expense. I told it this didn't sound like a big problem and asked for real cases. The cases showed moderate costs and rulings leaning against disallowance, so the brief now says "moderately costly, unknown error rate".
- **Switching ideas.** I asked for other topics and got five, ranked above TDS on impact. I stayed with TDS because of my time limit and the clearer split between AI and rules.

What came from me.

- Questioning whether this only saves accountants' time, which pushed us past an efficiency pitch.
- The idea that the tool lets a less qualified accountant do this work. When we examined it, this produced the weak-reviewer risk and the split between checking facts and making legal calls.
- Asking why TDS exists and comparing it with the US, which became the jurisdiction note.
- My first product idea was a prompt-based assistant. It lost to an automatic check in discussion, and I list it under rejected alternatives.

I did not get practitioner input, and I say so above.

One thing I'd defend without AI. The split between exact and probabilistic work: the LLM reads text and suggests a category, code does every rate and amount, and a person decides. I can explain why that boundary sits there and what breaks if it moves.